

# TinyML-Based Adaptive Speed Control for Car Robot: A Comparative Approach

Student: Alexandru Petriceanu<sup>1</sup> | Responsible Professor: Qing Wang<sup>1</sup>  
<sup>1</sup>EEMCS, Delft University of Technology, Netherlands

## 1. Background

- Adaptive speed control** heavily relies on how a navigator perceives their surroundings. Previously relied on ground truth systems like:
  - Stereo vision [1]
  - LiDAR [2]
  - Radar [3]
- Problems:** Computationally and energy inefficient. Infeasible for large-scale deployment.
- Solution:** ML-powered depth estimation model.
  - Use a U-Net or pyramid architecture model** to predict depths from images, using monocular vision [4, 5].
  - Computationally cheaper**, better fit for everyday use.
- State-of-the-art models** need GPUs and extensive resources to function [6, 7].
- Therefore, use **TinyML**, a subset of small models that run on cheap microcontrollers.
- For testing, use the Raspberry Pi Pico, which has only **264 KB of RAM** and **2 MB of flash memory**. Also uses the RP2040 chip, which has a **dual core Arm Cortex-M0+** processor.

## 2. Research Question

What is the post-compression efficiency of TinyML depth perception models when run on the Raspberry Pi Pico?

- Subquestions:**
  - What **other literature** is there on TinyML depth estimation?
  - Is running the **monocular depth estimation** task on the Raspberry Pi Pico **feasible**?
  - What effects do **compression** techniques such as **quantization** and **pruning** have on depth perception models?
  - What are representative **metrics** for **efficiency**?



Figure 1: Raspberry Pi Pico

## 3. Methodology



Figure 2.1: Before Preprocessing

- Three models** were selected and compared: **L-EfficientUNet** [4], **L-Enet** [8], **μPyD-Net** [5]
- One more original model was added to evaluate the efficiency of LSTMs: **Temporal-μPyD-Net**
- Models were trained and tested on the **Eigen split** of the **KITTI dataset** [9, 10]. This dataset contains over 60 recordings from stereo vision cameras and sparse LiDAR depth maps.

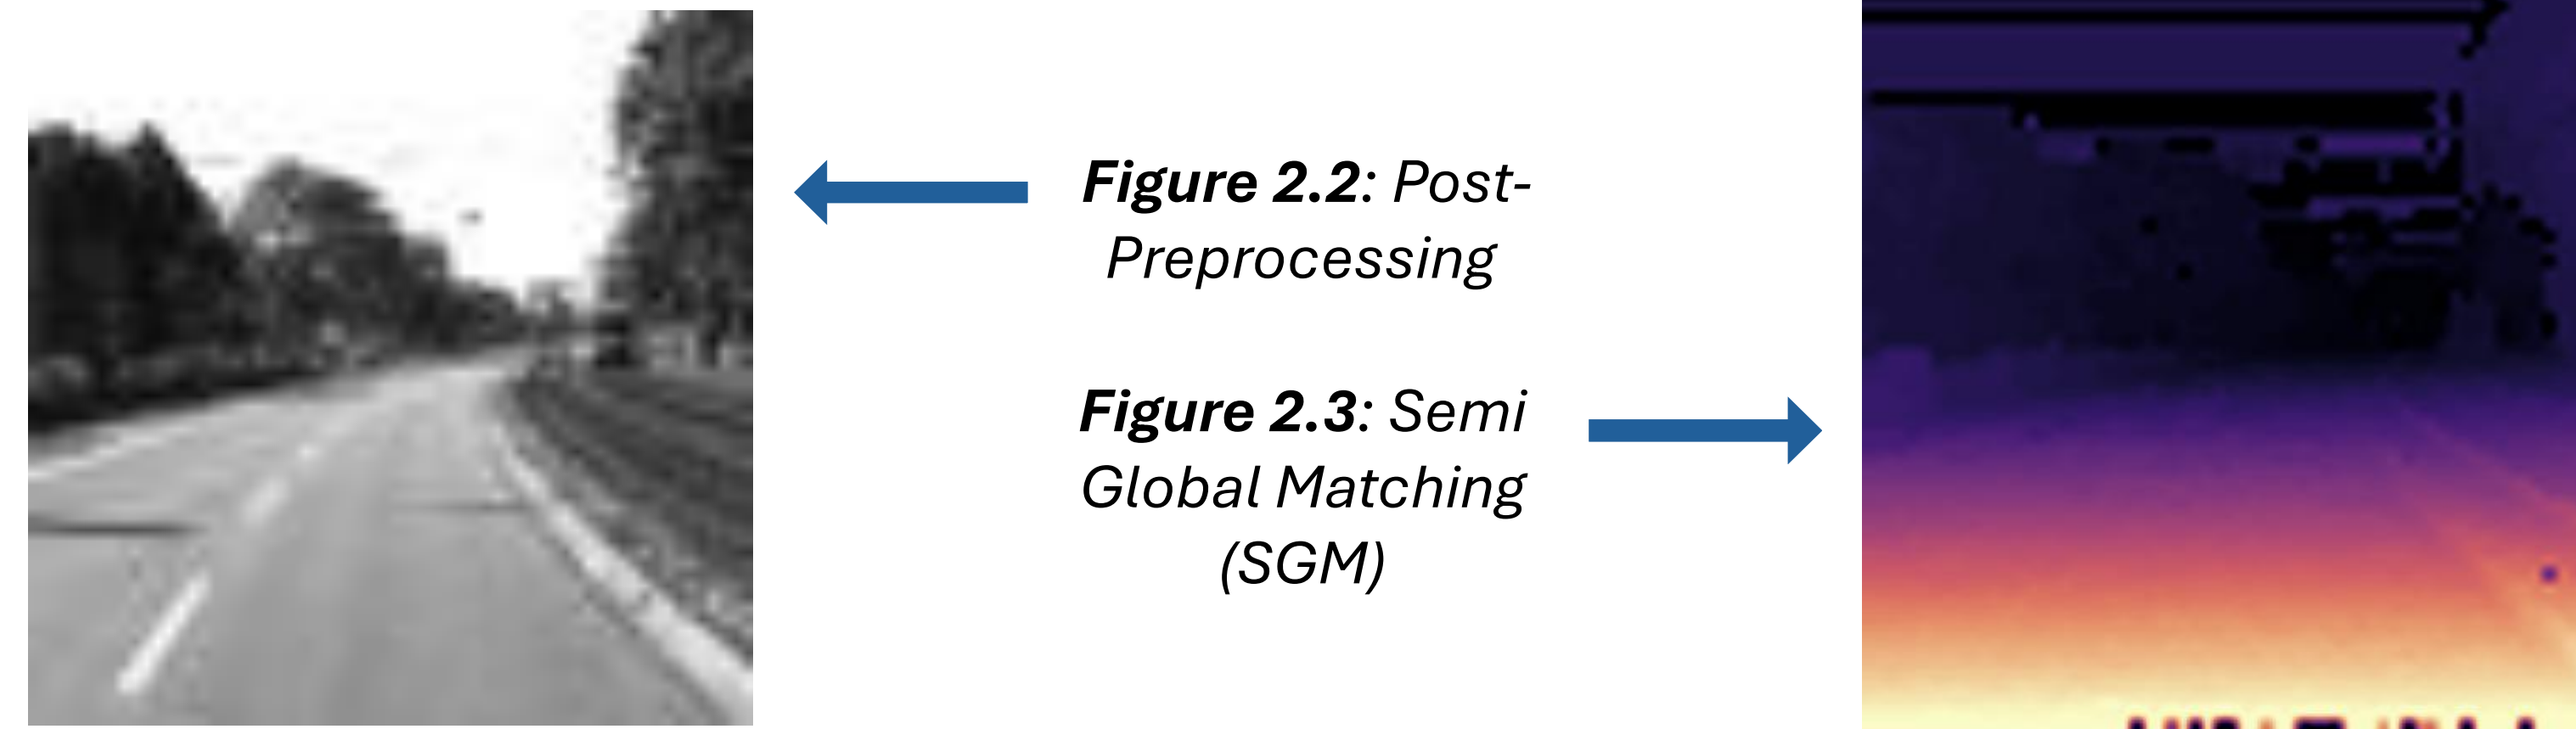


Figure 2.2: Post-Preprocessing

Figure 2.3: Semi Global Matching (SGM)

- The supervision label was the **SGM (Semi Global Matching) disparity map** obtained from each stereo image pair [11].
- Grayscale images from the left camera were cropped and resized to 32x32 pixels to be used as input.
- Models were trained using berHu (reverse Huber) loss and an Adam optimizer for 100 epochs.
- Evaluation** was performed using threshold accuracy, inference time, and memory used.
- The **top two best-performing** models were run for inference on the Raspberry Pi Pico.

## 4. Results

Model	$\delta < 1.25$	$\delta < 1.25^2$	$\delta < 1.25^3$
L-EfficientUNet	54.40%	70.08%	80.77%
L-Enet	55.88%	73.24%	83.35%
μPyD-Net	74.32%	83.95%	88.44%
Temporal-μPyD-Net	74.38%	83.68%	88.40%

Table 1: Accuracies for 64x64 resolution




Figure 3: μPyD-Net vs Temporal- μPyD-Net on 32x32 resolution




Figure 4: μPyD-Net vs Temporal- μPyD-Net SRAM and Inference Time (32x32)




Figure 5: Pre- vs Post-Quantization Accuracies




Figure 6: μPyD-Net Prediction




Figure 7: Temporal-μPyD-Net Prediction



Figure 8: μPyD-Net Prediction



Figure 9: Temporal-μPyD-Net Prediction



Figure 10: μPyD-Net Prediction (Overlaid)



Figure 11: Temporal-μPyD-Net Prediction (Overlaid)

## 5. Conclusion and Future Work

- Results show that running the depth perception task **is feasible** on the Raspberry Pi Pico. For **practical applications**, however, either more work on **fine-tuning inference time** or using a **better board** is recommended.
- Moreover, by measuring accuracies **pre-quantization** versus **post-quantization**, we can tell that full INT8 quantization does not affect accuracy in any meaningful way.
- Finally, we can conclude that **disparity maps** produced by the **SGM algorithm** are good supervision labels [11] and **berHu loss** facilitates good training results by not getting stuck in local minima and being a popular choice when it comes to depth estimation [5].